

1 !nvidia-smi

Sun Apr 10 23:08:23 2022

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI		460.32.03		Driver Version: 460.32.03			CUDA Version: 11.2		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU	Name	Persistence-M		Bus-Id	Disp.A		Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute M.	MIG M.	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
0	Tesla T4	Off		00000000:00:04.0	Off		0	0	
N/A	37C	P8	9W / 70W	0MiB / 15109MiB		0%	Default	N/A	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name		GPU Memory		
	ID	ID					Usage		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
No running processes found									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

1 !sudo apt-get install tree -qq > /dev/null

debconf: unable to initialize frontend: Dialog  
debconf: (No usable dialog-like program is installed, so the dialog based frontend cannot be used. at /usr/share/perl5/Debconf/FrontEnd/  
debconf: falling back to frontend: Readline  
debconf: unable to initialize frontend: Readline  
debconf: (This frontend requires a controlling tty.)  
debconf: falling back to frontend: Teletype  
dpkg-preconfigure: unable to re-open stdin:

```
1 #install libraries
2 !pip install torch==1.5.1+cu101 torchvision==0.6.1+cu101 -f https://download.pytorch.org/whl/torch_stable.html
3 !pip install numpy==1.17.3
4 !pip install PyYAML==5.3.1
5 !pip install git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI
```

```

Looking in links: https://download.pytorch.org/whl/torch\_stable.html
Collecting torch==1.5.1+cu101
  Downloading https://download.pytorch.org/whl/cu101/torch-1.5.1%2Bcu101-cp37-cp37m-linux\_x86\_64.whl
    704.4 MB 1.3 kB/s
Collecting torchvision==0.6.1+cu101
  Downloading https://download.pytorch.org/whl/cu101/torchvision-0.6.1%2Bcu101-cp37-cp37m-linux\_x86\_64.whl
    6.6 MB 40.6 MB/s
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torch)
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from torchvision)
Requirement already satisfied: pillow>4.1.1 in /usr/local/lib/python3.7/dist-packages (from torchvision)
Installing collected packages: torch, torchvision
  Attempting uninstall: torch
    Found existing installation: torch 1.10.0+cu111
    Uninstalling torch-1.10.0+cu111:
      Successfully uninstalled torch-1.10.0+cu111
  Attempting uninstall: torchvision
    Found existing installation: torchvision 0.11.1+cu111
    Uninstalling torchvision-0.11.1+cu111:
      Successfully uninstalled torchvision-0.11.1+cu111
ERROR: pip's dependency resolver does not currently take into account all the packages that are currently installed.
torchtext 0.11.0 requires torch==1.10.0, but you have torch 1.5.1+cu101 which is incompatible.
torchaudio 0.10.0+cu111 requires torch==1.10.0, but you have torch 1.5.1+cu101 which is incompatible.
Successfully installed torch-1.5.1+cu101 torchvision-0.6.1+cu101
Collecting numpy==1.17.3
  Downloading numpy-1.17.3-cp37-cp37m-manylinux1_x86_64.whl (20.0 MB)
    20.0 MB 31.4 MB/s
Installing collected packages: numpy
  Attempting uninstall: numpy
    Found existing installation: numpy 1.21.5
    Uninstalling numpy-1.21.5:
      Successfully uninstalled numpy-1.21.5
ERROR: pip's dependency resolver does not currently take into account all the packages that are currently installed.
tensorflow 2.8.0 requires tf-estimator-nightly==2.8.0.dev2021122109, which is not installed.
torchtext 0.11.0 requires torch==1.10.0, but you have torch 1.5.1+cu101 which is incompatible.
tensorflow 2.8.0 requires numpy>=1.20, but you have numpy 1.17.3 which is incompatible.
tables 3.7.0 requires numpy>=1.19.0, but you have numpy 1.17.3 which is incompatible.
kapre 0.3.7 requires numpy>=1.18.5, but you have numpy 1.17.3 which is incompatible.
jaxlib 0.3.2+cuda11.cudnn805 requires numpy>=1.19, but you have numpy 1.17.3 which is incompatible.
jax 0.3.4 requires numpy>=1.19, but you have numpy 1.17.3 which is incompatible.
datascience 0.10.6 requires folium==0.2.1, but you have folium 0.8.3 which is incompatible.
alumentations 0.1.12 requires imgaug<0.2.7,>=0.2.5, but you have imgaug 0.2.9 which is incompatible.
Successfully installed numpy-1.17.3
Collecting PyYAML==5.3.1
  Downloading PyYAML-5.3.1.tar.gz (269 kB)
    269 kB 29.4 MB/s
Building wheels for collected packages: PyYAML
  Building wheel for PyYAML (setup.py) ... done
  Created wheel for PyYAML: filename=PyYAML-5.3.1-cp37-cp37m-linux_x86_64.whl size=44636 sha256=...
  Stored in directory: /root/.cache/pip/wheels/5e/03/1e/e1e954795d6f35dfc7b637fe2277bfff
Successfully built PyYAML
Installing collected packages: PyYAML
  Attempting uninstall: PyYAML
    Found existing installation: PyYAML 3.13
    Uninstalling PyYAML-3.13:
      Successfully uninstalled PyYAML-3.13
Successfully installed PyYAML-5.3.1
Collecting git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI
  Cloning https://github.com/cocodataset/cocoapi.git to /tmp/pip-req-build-2ofmk09r
  Running command git clone -q https://github.com/cocodataset/cocoapi.git /tmp/pip-req-build-2ofmk09r
Requirement already satisfied: setuptools>=18.0 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: cython>=0.27.3 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: matplotlib>=2.1.0 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: cyclizer>=0.10 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI)
Building wheels for collected packages: pycocotools
  Building wheel for pycocotools (setup.py) ... done
  Created wheel for pycocotools: filename=pycocotools-2.0-cp37-cp37m-linux_x86_64.whl size=11111 sha256=...
  Stored in directory: /tmp/pip-ephem-wheel-cache-i872y92t/wheels/e2/6b/1d/344ac773c7495...
Successfully built pycocotools
Installing collected packages: pycocotools
  Attempting uninstall: pycocotools
    Found existing installation: pycocotools 2.0.4
    Uninstalling pycocotools-2.0.4:
      Successfully uninstalled pycocotools-2.0.4
Successfully installed pycocotools-2.0

```

## Build the Dataset

```
1 # imports
2 from pathlib import Path
3 from tqdm import tqdm
4 import numpy as np
5 import json
6 import urllib
7 import PIL.Image as Image
8 import cv2
9 import torch
10 import torchvision
11 from IPython.display import display
12 from sklearn.model_selection import train_test_split
13 import os
14 import PIL
15
16 import seaborn as sns
17 from pylab import rcParams
18 import matplotlib.pyplot as plt
19 from matplotlib import rc
20
21 %matplotlib inline
22 %config InlineBackend.figure_format='retina'
23 sns.set(style='whitegrid', palette='muted', font_scale=1.2)
24 rcParams['figure.figsize'] = 16, 10
25
26 np.random.seed(42)


1 # our data directories
2 !tree Weapon-detection-1 -L 2


    Weapon-detection-1 [error opening dir]


    0 directories, 0 files


1 !pip install roboflow
2
3 from roboflow import Roboflow
4 rf = Roboflow(api_key="HAAZzUmIip5OvdPutJ90")
5 project = rf.workspace("alon-barak").project("weapon-detection-dp9bo")
6 dataset = project.version(1).download("yolov5")
7


1 # cloning the git repo
2 !git clone https://github.com/ultralytics/yolov5
3 %cd yolov5
4 !git checkout 741fac815e366d74eed020efb8c68a23828ee3e9


1 # show the dimensions of the images in the dataset
2 img = cv2.imread('/content/yolov5/Weapon-detection-1/train/images/107_jpeg.rf.22d3c90341d7af62ed4cb51f334ce3dc.jpg', cv2.IMREAD_UNCHANGED)
3 dimensions = img.shape
4 # height, width, number of channels in image
5 height = img.shape[0]
6 width = img.shape[1]
7 channels = img.shape[2]
8
9 print('Image Dimension    : ',dimensions)
10 print('Image Height       : ',height)
11 print('Image Width          : ',width)
12 print('Number of Channels    : ',channels)


    Image Dimension      : (460, 460, 3)
    Image Height         : 460
    Image Width          : 460
    Number of Channels   : 3


1 img = cv2.imread('/content/yolov5/Weapon-detection-1/train/images/107_jpeg.rf.22d3c90341d7af62ed4cb51f334ce3dc.jpg', cv2.IMREAD_UNCHANGED)
2 plt.imshow(img)
```


A 5x5 grid of 25 frames showing a person in a dark environment, possibly a laboratory or industrial setting, with a bright light source visible in the upper right corner. The person is wearing a dark jacket and a blue mask. The frames are arranged in a 5x5 grid, with the top-left frame being the most visible and the bottom-right frame being the least visible. The grid is overlaid on a dark background, and the person is positioned in the center-left of the frame. The bright light source is located in the upper right corner, creating a strong contrast with the dark environment. The overall image is dark and grainy, with a high level of contrast between the light source and the surrounding environment.

```
1 !python train.py --img 480 --batch 4 --epochs 30 --data ./data/data.yaml --cfg ./models/yolov5x.yaml --weights yolov5x.pt --name yolov5x,
```

**hyperparameters:** lr=0.01, lr\_f=0.01, momentum=0.937, weight\_decay=0.0005, warmup\_epochs=3.0, warmup\_momentum=0.8, warmup\_bias\_lr=0.1,  
**Weights & Biases:** run 'pip install wandb' to automatically track and visualize YOLOv5 🚀 runs (RECOMMENDED)  
**TensorBoard:** Start with 'tensorboard --logdir runs/train', view at <http://localhost:6006/>

```
Transferred 738/745 items from yolov5x.pt
Scaled weight_decay = 0.0005
optimizer: SGD with parameter groups 123 weight (no decay), 126 weight, 126 bias
```

**albumentations:** version 1.0.3 required by YOLOv5, but version 0.1.12 is currently installed  
**train:** Scanning '/content/yolov5/Weapon-detection-1/train/labels' images and labels...437 found, 0 missing, 0 empty, 0 corrupt: 100%  
**train:** New cache created: /content/yolov5/Weapon-detection-1/train/labels.cache  
**val:** Scanning '/content/yolov5/Weapon-detection-1/valid/labels' images and labels...30 found, 0 missing, 0 empty, 0 corrupt: 100% 30/  
**val:** New cache created: /content/yolov5/Weapon-detection-1/valid/labels.cache  
 Plotting labels to runs/train/yolov5x\_weapon2/labels.jpg...

**AutoAnchor:** 4.61 anchors/target, 1.000 Best Possible Recall (BPR). Current anchors are a good fit to dataset   
 Image sizes 480 train, 480 val  
 Using 2 dataloader workers  
 Logging results to runs/train/yolov5x\_weapon2  
 Starting training for 30 epochs...

Epoch	gpu_mem	box	obj	cls	labels	img_size						
0/29	3.4G	0.09619	0.02528	0	3	480: 100% 110/110	[00:28<00:00,	3.87it/s]				
	Class	Images	Labels	P	R	mAP@.5	mAP@.5:.95:	100% 4/4	[00:00<00:00,	6.75it/s]		
	all	30	47	0.218	0.383	0.139	0.039					

```

1 # view some predictions
2 img = cv2.imread('/content/yolov5/runs/train/yolov5x_weapon2/val_batch2_pred.jpg', cv2.IMREAD_UNCHANGED)
3 plt.imshow(img)

```

<matplotlib.image.AxesImage at 0x7fcd9f7fbc90>

